

Data Mining Final Project

- Rajendra Prasad Patil

Glossary:

- Import libraries
- Load dataset
- Analysis on dataset
- Splitting the dataset into labels and features
- Performing normalization on dataset
- Splitting dataset using K fold
- Running the model
 - SVM Model
 - K Nearest Neighbors
 - Random Forest Classifier
- Output Performance Metrics

Importing libraries

In [1]:

```
import pandas as pd
import numpy as np
from sklearn.datasets import load_breast_cancer
from sklearn.preprocessing import StandardScaler

import warnings
warnings.filterwarnings('ignore')

# SVM classifier
from sklearn import svm

# KNN classifier
from sklearn.neighbors import KNeighborsClassifier

# Import Random Forest Model
from sklearn.ensemble import RandomForestClassifier

# Import libraries for lstm classification
from keras.layers import Dense, Dropout, LSTM, Embedding
from keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential

# for checking the model accuracy
from sklearn.metrics import confusion_matrix, accuracy_score, classification_
```

Loading the dataset

```
In [2]: dataset = load_breast_cancer()
```

```
In [3]: input_length = len(dataset['data'][0])
```

Preliminary analysis

```
In [4]: class_names = dataset['target_names']
print('Target variables : ', class_names)

(unique, counts) = np.unique(dataset['target'], return_counts=True)

print('Unique values of the target variable', unique)
print('Counts of the target variable : ', counts)
```

```
Target variables : ['malignant' 'benign']
Unique values of the target variable [0 1]
Counts of the target variable : [212 357]
```

- The dataset is suited for binary classification
- The dataset has no skewed nature

The data is split into features and labels

```
In [5]: X = dataset['data']
y = dataset['target']
```

Apply normalization operation for numerical stability

```
In [6]: standardizer = StandardScaler()
X = standardizer.fit_transform(X)
```

Performance Metrics

Function to calculate all the available performance metrics

In [7]:

```

performance_metrics = ['True Negative', 'False Positive', 'False Negative', '
                        'Precision', 'Accuracy', 'F1 Score', 'Error Rate', 'Ne
                        'False Discovery Rate', 'False Negative Rate', 'Balanc
                        'Heidke Skill Score']

def compute_performance_metrics(prediction, y_test, df, is_lstm = False):

    if is_lstm:
        threshold = 0.80
        for i, each in enumerate(prediction):
            if each[0] > threshold:
                prediction[i] = 1
            else:
                prediction[i] = 0

    TN, FP, FN, TP = confusion_matrix(y_test, prediction).ravel()

    sensitivity = TP / (TP + FN)
    specificity = TN / (FP + TN)
    precision = TP / (TP + FP)
    accuracy = (TP+TN) / (TP+FP+TN+FN)
    f1_score = 2 * TP / ((2 * TP) + FP + FN)
    error_rate = (FP + FN) / (TP + FP + FN + TN)
    negative_predicted_value = TN / (TN + FN)
    false_positive_rate = FP / (FP + TN)
    false_discovery_rate = FP / (FP + TP)
    false_negative_rate = FN / (FN + TP)
    balanced_accuracy = 0.5 * ((TP / (TP + FN)) + (TN / (TN + FP)))
    true_skill_statistics = ((TP / (TP + FN)) - (FP / (TN + FP)))
    heidke_skill_score = 2 * ((TP * TN) - (FP * FN)) / (((TP + FN) * (FN + T

    df = df.append({performance_metrics[0]: TN, performance_metrics[1]: FP, p
                    performance_metrics[3]: TP, performance_metrics[4]: sensi
                    performance_metrics[6]: precision, performance_metrics[7]
                    performance_metrics[9]: error_rate, performance_metrics[1
                    performance_metrics[11]: false_positive_rate, performance
                    performance_metrics[13]: false_negative_rate, performance
                    balanced_accuracy, performance_metrics[15]: true_skill_st
                    performance_metrics[16]: heidke_skill_score}, ignore_inde

    return df

```

K-fold cross validation

In [8]:

```

from sklearn.model_selection import KFold
kfold = KFold(n_splits=10, shuffle=True, random_state=0)

```

Dataframes for performance metrics

In [9]:

```

svm_metrics_df = pd.DataFrame(columns=performance_metrics)
kn_metrics_df = pd.DataFrame(columns=performance_metrics)
rf_metrics_df = pd.DataFrame(columns=performance_metrics)
lstm_metrics_df = pd.DataFrame(columns=performance_metrics)

```

SVM Model

In [10]:

```
svm_model = svm.SVC()
for train_index, test_index in kfold.split(X):
    X_train, X_test, y_train, y_test = X[train_index], X[test_index], y[train_index], y[test_index]

    # we train the algorithm with training data and training output
    svm_model.fit(X_train, y_train)

    # we pass the testing data to the stored algorithm to predict the outcome
    prediction = svm_model.predict(X_test)

    # print metrics
    svm_metrics_df = compute_performance_metrics(prediction, y_test, svm_model)

svm_metrics_df.index += 1
svm_metrics_df.loc['Average'] = svm_metrics_df.mean()
```

In [11]:

```
svm_metrics_df
```

Out[11]:

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy |
|---------|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|
| 1 | 22.0 | 0.0 | 0.0 | 35.0 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |
| 2 | 23.0 | 2.0 | 1.0 | 31.0 | 0.968750 | 0.920000 | 0.939394 | 0.947368 |
| 3 | 15.0 | 1.0 | 1.0 | 40.0 | 0.975610 | 0.937500 | 0.975610 | 0.964912 |
| 4 | 20.0 | 0.0 | 0.0 | 37.0 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |
| 5 | 17.0 | 1.0 | 0.0 | 39.0 | 1.000000 | 0.944444 | 0.975000 | 0.982456 |
| 6 | 19.0 | 3.0 | 0.0 | 35.0 | 1.000000 | 0.863636 | 0.921053 | 0.947368 |
| 7 | 22.0 | 1.0 | 1.0 | 33.0 | 0.970588 | 0.956522 | 0.970588 | 0.964912 |
| 8 | 23.0 | 0.0 | 2.0 | 32.0 | 0.941176 | 1.000000 | 1.000000 | 0.964912 |
| 9 | 18.0 | 0.0 | 0.0 | 39.0 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |
| 10 | 25.0 | 0.0 | 0.0 | 31.0 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |
| Average | 20.4 | 0.8 | 0.5 | 35.2 | 0.985612 | 0.962210 | 0.978164 | 0.977193 |

K-Nearest Neighbors

```
In [12]: model = KNeighborsClassifier(n_neighbors=3) # this examines 3 neighbors for p
for train_index, test_index in kfold.split(X):
    X_train, X_test, y_train, y_test = X[train_index], X[test_index], y[train_index], y[test_index]

    # we train the algorithm with training data and training output
    model.fit(X_train, y_train)

    # we pass the testing data to the stored algorithm to predict the outcome
    prediction = model.predict(X_test)

    # print metrics
    kn_metrics_df = compute_performance_metrics(prediction, y_test, kn_metrics_df)

kn_metrics_df.index += 1
kn_metrics_df.loc['Average'] = kn_metrics_df.mean()
```

```
In [13]: kn_metrics_df
```

```
Out[13]:
```

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy |
|---------|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|
| 1 | 21.0 | 1.0 | 0.0 | 35.0 | 1.000000 | 0.954545 | 0.972222 | 0.982456 |
| 2 | 21.0 | 4.0 | 0.0 | 32.0 | 1.000000 | 0.840000 | 0.888889 | 0.929825 |
| 3 | 14.0 | 2.0 | 1.0 | 40.0 | 0.975610 | 0.875000 | 0.952381 | 0.947368 |
| 4 | 19.0 | 1.0 | 0.0 | 37.0 | 1.000000 | 0.950000 | 0.973684 | 0.982456 |
| 5 | 17.0 | 1.0 | 1.0 | 38.0 | 0.974359 | 0.944444 | 0.974359 | 0.964912 |
| 6 | 19.0 | 3.0 | 0.0 | 35.0 | 1.000000 | 0.863636 | 0.921053 | 0.947368 |
| 7 | 22.0 | 1.0 | 0.0 | 34.0 | 1.000000 | 0.956522 | 0.971429 | 0.982456 |
| 8 | 23.0 | 0.0 | 0.0 | 34.0 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |
| 9 | 18.0 | 0.0 | 0.0 | 39.0 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |
| 10 | 23.0 | 2.0 | 0.0 | 31.0 | 1.000000 | 0.920000 | 0.939394 | 0.964286 |
| Average | 19.7 | 1.5 | 0.2 | 35.5 | 0.994997 | 0.930415 | 0.959341 | 0.970113 |

Random Forest Classifier

In [14]:

```
#Create a Gaussian Classifier
model = RandomForestClassifier(n_estimators=100)

for train_index, test_index in kfold.split(X):
    X_train, X_test, y_train, y_test = X[train_index], X[test_index], y[train_index], y[test_index]

    # we train the algorithm with training data and training output
    model.fit(X_train, y_train)

    # we pass the testing data to the stored algorithm to predict the outcome
    prediction = model.predict(X_test)

    # print metrics
    rf_metrics_df = compute_performance_metrics(prediction, y_test, rf_metrics_df)

rf_metrics_df.index += 1
rf_metrics_df.loc['Average'] = rf_metrics_df.mean()
```

In [15]:

```
rf_metrics_df
```

Out[15]:

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy |
|---------|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|
| 1 | 21.0 | 1.0 | 2.0 | 33.0 | 0.942857 | 0.954545 | 0.970588 | 0.947368 |
| 2 | 24.0 | 1.0 | 0.0 | 32.0 | 1.000000 | 0.960000 | 0.969697 | 0.982456 |
| 3 | 15.0 | 1.0 | 2.0 | 39.0 | 0.951220 | 0.937500 | 0.975000 | 0.947368 |
| 4 | 19.0 | 1.0 | 0.0 | 37.0 | 1.000000 | 0.950000 | 0.973684 | 0.982456 |
| 5 | 17.0 | 1.0 | 1.0 | 38.0 | 0.974359 | 0.944444 | 0.974359 | 0.964912 |
| 6 | 17.0 | 5.0 | 2.0 | 33.0 | 0.942857 | 0.772727 | 0.868421 | 0.877193 |
| 7 | 22.0 | 1.0 | 0.0 | 34.0 | 1.000000 | 0.956522 | 0.971429 | 0.982456 |
| 8 | 20.0 | 3.0 | 1.0 | 33.0 | 0.970588 | 0.869565 | 0.916667 | 0.929825 |
| 9 | 18.0 | 0.0 | 0.0 | 39.0 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |
| 10 | 24.0 | 1.0 | 0.0 | 31.0 | 1.000000 | 0.960000 | 0.968750 | 0.982143 |
| Average | 19.7 | 1.5 | 0.8 | 34.9 | 0.978188 | 0.930530 | 0.958859 | 0.959618 |

LSTM classifier

In [16]:

```

model = Sequential()
model.add(LSTM(20, input_shape=(input_length, 1)))
model.add(Dropout(0.5))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

for train_index, test_index in kfold.split(X):
    print('*'*100)
    X_train, X_test, y_train, y_test = X[train_index], X[test_index], y[train_index], y[test_index]

    # we train the algorithm with training data and training output
    model.fit(X_train, y_train, batch_size=8, epochs=10, validation_data=(X_test, y_test))

    # we pass the testing data to the stored algorithm to predict the outcome
    prediction = model.predict(X_test)

    # print metrics
    lstm_metrics_df = compute_performance_metrics(prediction, y_test, lstm_metrics_df)

lstm_metrics_df.index += 1
lstm_metrics_df.loc['Average'] = lstm_metrics_df.mean()

```

Metal device set to: Apple M1

2021-12-04 16:56:43.997537: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:305] Could not identify NUMA node of platform GPU ID 0, defaulting to 0. Your kernel may not have been built with NUMA support.

2021-12-04 16:56:43.997965: I tensorflow/core/common_runtime/pluggable_device/pluggable_device_factory.cc:271] Created TensorFlow device (/job:localhost/replica:0/task:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevice (device: 0, name: METAL, pci bus id: <undefined>)

2021-12-04 16:56:44.209424: W tensorflow/core/platform/profile_utils/cpu_utils.cc:128] Failed to get CPU frequency: 0 Hz

Epoch 1/10

2021-12-04 16:56:44.590973: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.

2021-12-04 16:56:44.711991: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.

5/64 [=>.....] - ETA: 0s - loss: 0.8705 - accuracy: 0.2250

2021-12-04 16:56:44.831839: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.

64/64 [=====] - 2s 18ms/step - loss: 0.7669 - accuracy: 0.3457 - val_loss: 0.7148 - val_accuracy: 0.4912

2021-12-04 16:56:45.915998: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.

2021-12-04 16:56:45.966215: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.

Epoch 2/10

64/64 [=====] - 1s 15ms/step - loss: 0.6055 - accuracy: 0.7324 - val_loss: 0.4008 - val_accuracy: 0.9298

Epoch 3/10

```
64/64 [=====] - 1s 15ms/step - loss: 0.3536 - accuracy: 0.9043 - val_loss: 0.3022 - val_accuracy: 0.8947
Epoch 4/10
64/64 [=====] - 1s 14ms/step - loss: 0.3118 - accuracy: 0.9199 - val_loss: 0.2717 - val_accuracy: 0.9123
Epoch 5/10
64/64 [=====] - 1s 15ms/step - loss: 0.2680 - accuracy: 0.9277 - val_loss: 0.2278 - val_accuracy: 0.9474
Epoch 6/10
64/64 [=====] - 1s 15ms/step - loss: 0.2960 - accuracy: 0.9238 - val_loss: 0.2116 - val_accuracy: 0.9474
Epoch 7/10
64/64 [=====] - 1s 15ms/step - loss: 0.2694 - accuracy: 0.9219 - val_loss: 0.1653 - val_accuracy: 0.9649
Epoch 8/10
64/64 [=====] - 1s 15ms/step - loss: 0.2853 - accuracy: 0.9180 - val_loss: 0.2820 - val_accuracy: 0.9123
Epoch 9/10
64/64 [=====] - 1s 15ms/step - loss: 0.3803 - accuracy: 0.8848 - val_loss: 0.2712 - val_accuracy: 0.9123
Epoch 10/10
64/64 [=====] - 1s 15ms/step - loss: 0.3596 - accuracy: 0.8789 - val_loss: 0.2642 - val_accuracy: 0.9123
*****
*****
Epoch 1/10
5/64 [=>.....] - ETA: 0s - loss: 0.3474 - accuracy: 0.8750
2021-12-04 16:56:54.706389: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
2021-12-04 16:56:54.742269: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
64/64 [=====] - 1s 15ms/step - loss: 0.3186 - accuracy: 0.8906 - val_loss: 0.3013 - val_accuracy: 0.8947
Epoch 2/10
64/64 [=====] - 1s 14ms/step - loss: 0.3233 - accuracy: 0.8945 - val_loss: 0.3014 - val_accuracy: 0.8947
Epoch 3/10
64/64 [=====] - 1s 15ms/step - loss: 0.3109 - accuracy: 0.8926 - val_loss: 0.2979 - val_accuracy: 0.8947
Epoch 4/10
64/64 [=====] - 1s 15ms/step - loss: 0.2944 - accuracy: 0.9004 - val_loss: 0.2953 - val_accuracy: 0.8947
Epoch 5/10
64/64 [=====] - 1s 14ms/step - loss: 0.2919 - accuracy: 0.9062 - val_loss: 0.3285 - val_accuracy: 0.8947
Epoch 6/10
64/64 [=====] - 1s 15ms/step - loss: 0.2684 - accuracy: 0.9121 - val_loss: 0.2973 - val_accuracy: 0.8947
Epoch 7/10
64/64 [=====] - 1s 15ms/step - loss: 0.2754 - accuracy: 0.9082 - val_loss: 0.2935 - val_accuracy: 0.8947
Epoch 8/10
64/64 [=====] - 1s 15ms/step - loss: 0.2699 - accuracy: 0.9102 - val_loss: 0.3311 - val_accuracy: 0.8947
Epoch 9/10
64/64 [=====] - 1s 15ms/step - loss: 0.3057 - accuracy: 0.9102 - val_loss: 0.4018 - val_accuracy: 0.8421
Epoch 10/10
```



```
64/64 [=====] - 1s 15ms/step - loss: 0.2662 - accuracy: 0.9180 - val_loss: 0.3247 - val_accuracy: 0.8596
*****
*****
Epoch 1/10
64/64 [=====] - 1s 15ms/step - loss: 0.2479 - accuracy: 0.9141 - val_loss: 0.3898 - val_accuracy: 0.8596
Epoch 2/10
64/64 [=====] - 1s 14ms/step - loss: 0.2531 - accuracy: 0.9141 - val_loss: 0.3712 - val_accuracy: 0.8596
Epoch 3/10
64/64 [=====] - 1s 14ms/step - loss: 0.2453 - accuracy: 0.9121 - val_loss: 0.2918 - val_accuracy: 0.8947
Epoch 4/10
64/64 [=====] - 1s 15ms/step - loss: 0.2315 - accuracy: 0.9160 - val_loss: 0.2987 - val_accuracy: 0.8772
Epoch 5/10
64/64 [=====] - 1s 15ms/step - loss: 0.2359 - accuracy: 0.9160 - val_loss: 0.2014 - val_accuracy: 0.9298
Epoch 6/10
64/64 [=====] - 1s 14ms/step - loss: 0.2201 - accuracy: 0.9160 - val_loss: 0.2331 - val_accuracy: 0.9123
Epoch 7/10
64/64 [=====] - 1s 15ms/step - loss: 0.2199 - accuracy: 0.9141 - val_loss: 0.2189 - val_accuracy: 0.9123
Epoch 8/10
64/64 [=====] - 1s 15ms/step - loss: 0.2314 - accuracy: 0.9141 - val_loss: 0.2772 - val_accuracy: 0.8947
Epoch 9/10
64/64 [=====] - 1s 15ms/step - loss: 0.2394 - accuracy: 0.9043 - val_loss: 0.2581 - val_accuracy: 0.8772
Epoch 10/10
64/64 [=====] - 1s 14ms/step - loss: 0.2241 - accuracy: 0.9082 - val_loss: 0.2674 - val_accuracy: 0.8947
*****
*****
Epoch 1/10
64/64 [=====] - 1s 15ms/step - loss: 0.2142 - accuracy: 0.9141 - val_loss: 0.1580 - val_accuracy: 0.9474
Epoch 2/10
64/64 [=====] - 1s 15ms/step - loss: 0.2356 - accuracy: 0.8926 - val_loss: 0.2279 - val_accuracy: 0.9123
Epoch 3/10
64/64 [=====] - 1s 15ms/step - loss: 0.2346 - accuracy: 0.9102 - val_loss: 0.2239 - val_accuracy: 0.9123
Epoch 4/10
64/64 [=====] - 1s 15ms/step - loss: 0.2043 - accuracy: 0.9160 - val_loss: 0.1522 - val_accuracy: 0.9474
Epoch 5/10
64/64 [=====] - 1s 14ms/step - loss: 0.2003 - accuracy: 0.9277 - val_loss: 0.1623 - val_accuracy: 0.9474
Epoch 6/10
64/64 [=====] - 1s 15ms/step - loss: 0.2188 - accuracy: 0.9199 - val_loss: 0.2400 - val_accuracy: 0.9123
Epoch 7/10
64/64 [=====] - 1s 15ms/step - loss: 0.2053 - accuracy: 0.9199 - val_loss: 0.2312 - val_accuracy: 0.9298
Epoch 8/10
64/64 [=====] - 1s 15ms/step - loss: 0.2161 - accuracy:
```

```
cy: 0.9141 - val_loss: 0.2019 - val_accuracy: 0.9298
Epoch 9/10
64/64 [=====] - 1s 15ms/step - loss: 0.2124 - accuracy: 0.9238 - val_loss: 0.1460 - val_accuracy: 0.9649
Epoch 10/10
64/64 [=====] - 1s 15ms/step - loss: 0.2222 - accuracy: 0.9082 - val_loss: 0.2161 - val_accuracy: 0.9298
*****
*****
Epoch 1/10
64/64 [=====] - 1s 16ms/step - loss: 0.2019 - accuracy: 0.9121 - val_loss: 0.2104 - val_accuracy: 0.9298
Epoch 2/10
64/64 [=====] - 1s 15ms/step - loss: 0.1840 - accuracy: 0.9277 - val_loss: 0.2069 - val_accuracy: 0.9474
Epoch 3/10
64/64 [=====] - 1s 15ms/step - loss: 0.1915 - accuracy: 0.9219 - val_loss: 0.1769 - val_accuracy: 0.9649
Epoch 4/10
64/64 [=====] - 1s 15ms/step - loss: 0.1822 - accuracy: 0.9238 - val_loss: 0.1090 - val_accuracy: 0.9649
Epoch 5/10
64/64 [=====] - 1s 16ms/step - loss: 0.1923 - accuracy: 0.9238 - val_loss: 0.1216 - val_accuracy: 0.9649
Epoch 6/10
64/64 [=====] - 1s 15ms/step - loss: 0.1908 - accuracy: 0.9316 - val_loss: 0.1075 - val_accuracy: 0.9649
Epoch 7/10
64/64 [=====] - 1s 15ms/step - loss: 0.1693 - accuracy: 0.9395 - val_loss: 0.1110 - val_accuracy: 0.9649
Epoch 8/10
64/64 [=====] - 1s 15ms/step - loss: 0.1678 - accuracy: 0.9316 - val_loss: 0.1257 - val_accuracy: 0.9825
Epoch 9/10
64/64 [=====] - 1s 15ms/step - loss: 0.1720 - accuracy: 0.9375 - val_loss: 0.1286 - val_accuracy: 0.9474
Epoch 10/10
64/64 [=====] - 1s 14ms/step - loss: 0.1693 - accuracy: 0.9414 - val_loss: 0.1326 - val_accuracy: 0.9474
*****
*****
Epoch 1/10
64/64 [=====] - 1s 15ms/step - loss: 0.1635 - accuracy: 0.9395 - val_loss: 0.3626 - val_accuracy: 0.8772
Epoch 2/10
64/64 [=====] - 1s 15ms/step - loss: 0.1831 - accuracy: 0.9277 - val_loss: 0.3327 - val_accuracy: 0.8947
Epoch 3/10
64/64 [=====] - 1s 15ms/step - loss: 0.1677 - accuracy: 0.9316 - val_loss: 0.2351 - val_accuracy: 0.9123
Epoch 4/10
64/64 [=====] - 1s 14ms/step - loss: 0.1603 - accuracy: 0.9336 - val_loss: 0.2359 - val_accuracy: 0.9474
Epoch 5/10
64/64 [=====] - 1s 15ms/step - loss: 0.1675 - accuracy: 0.9316 - val_loss: 0.3012 - val_accuracy: 0.9298
Epoch 6/10
64/64 [=====] - 1s 15ms/step - loss: 0.1513 - accuracy: 0.9395 - val_loss: 0.2920 - val_accuracy: 0.9298
```

```
Epoch 7/10
64/64 [=====] - 1s 15ms/step - loss: 0.1738 - accuracy: 0.9277 - val_loss: 0.2759 - val_accuracy: 0.9298
Epoch 8/10
64/64 [=====] - 1s 15ms/step - loss: 0.1522 - accuracy: 0.9453 - val_loss: 0.2827 - val_accuracy: 0.9298
Epoch 9/10
64/64 [=====] - 1s 14ms/step - loss: 0.1707 - accuracy: 0.9316 - val_loss: 0.2190 - val_accuracy: 0.9649
Epoch 10/10
64/64 [=====] - 1s 15ms/step - loss: 0.1744 - accuracy: 0.9180 - val_loss: 0.2344 - val_accuracy: 0.9474
*****
*****
Epoch 1/10
64/64 [=====] - 1s 15ms/step - loss: 0.1676 - accuracy: 0.9258 - val_loss: 0.2037 - val_accuracy: 0.9298
Epoch 2/10
64/64 [=====] - 1s 15ms/step - loss: 0.1861 - accuracy: 0.9258 - val_loss: 0.2001 - val_accuracy: 0.9474
Epoch 3/10
64/64 [=====] - 1s 15ms/step - loss: 0.1627 - accuracy: 0.9355 - val_loss: 0.2029 - val_accuracy: 0.9298
Epoch 4/10
64/64 [=====] - 1s 15ms/step - loss: 0.1667 - accuracy: 0.9395 - val_loss: 0.2240 - val_accuracy: 0.9123
Epoch 5/10
64/64 [=====] - 1s 15ms/step - loss: 0.1604 - accuracy: 0.9434 - val_loss: 0.2157 - val_accuracy: 0.9123
Epoch 6/10
64/64 [=====] - 1s 15ms/step - loss: 0.1699 - accuracy: 0.9414 - val_loss: 0.2422 - val_accuracy: 0.9123
Epoch 7/10
64/64 [=====] - 1s 15ms/step - loss: 0.1987 - accuracy: 0.9316 - val_loss: 0.2378 - val_accuracy: 0.9123
Epoch 8/10
64/64 [=====] - 1s 15ms/step - loss: 0.1646 - accuracy: 0.9492 - val_loss: 0.2165 - val_accuracy: 0.9123
Epoch 9/10
64/64 [=====] - 1s 14ms/step - loss: 0.1866 - accuracy: 0.9375 - val_loss: 0.2758 - val_accuracy: 0.9123
Epoch 10/10
64/64 [=====] - 1s 15ms/step - loss: 0.2179 - accuracy: 0.9375 - val_loss: 0.2483 - val_accuracy: 0.8947
*****
*****
Epoch 1/10
64/64 [=====] - 1s 15ms/step - loss: 0.2129 - accuracy: 0.9395 - val_loss: 0.1529 - val_accuracy: 0.9474
Epoch 2/10
64/64 [=====] - 1s 15ms/step - loss: 0.2066 - accuracy: 0.9414 - val_loss: 0.1470 - val_accuracy: 0.9474
Epoch 3/10
64/64 [=====] - 1s 15ms/step - loss: 0.2414 - accuracy: 0.9277 - val_loss: 0.1848 - val_accuracy: 0.9298
Epoch 4/10
64/64 [=====] - 1s 15ms/step - loss: 0.2199 - accuracy: 0.9297 - val_loss: 0.1057 - val_accuracy: 0.9825
Epoch 5/10
```

```

64/64 [=====] - 1s 15ms/step - loss: 0.1854 - accuracy: 0.9355 - val_loss: 0.1057 - val_accuracy: 0.9825
Epoch 6/10
64/64 [=====] - 1s 15ms/step - loss: 0.1809 - accuracy: 0.9551 - val_loss: 0.1102 - val_accuracy: 0.9649
Epoch 7/10
64/64 [=====] - 1s 15ms/step - loss: 0.2096 - accuracy: 0.9336 - val_loss: 0.1125 - val_accuracy: 0.9649
Epoch 8/10
64/64 [=====] - 1s 15ms/step - loss: 0.1885 - accuracy: 0.9414 - val_loss: 0.1413 - val_accuracy: 0.9649
Epoch 9/10
64/64 [=====] - 1s 15ms/step - loss: 0.1840 - accuracy: 0.9492 - val_loss: 0.0939 - val_accuracy: 0.9649
Epoch 10/10
64/64 [=====] - 1s 15ms/step - loss: 0.1776 - accuracy: 0.9492 - val_loss: 0.0966 - val_accuracy: 0.9649
*****
*****
Epoch 1/10
64/64 [=====] - 1s 15ms/step - loss: 0.1588 - accuracy: 0.9492 - val_loss: 0.0905 - val_accuracy: 0.9649
Epoch 2/10
64/64 [=====] - 1s 15ms/step - loss: 0.1563 - accuracy: 0.9512 - val_loss: 0.0881 - val_accuracy: 0.9474
Epoch 3/10
64/64 [=====] - 1s 15ms/step - loss: 0.1671 - accuracy: 0.9492 - val_loss: 0.1140 - val_accuracy: 0.9474
Epoch 4/10
64/64 [=====] - 1s 15ms/step - loss: 0.2109 - accuracy: 0.9238 - val_loss: 0.1048 - val_accuracy: 0.9649
Epoch 5/10
64/64 [=====] - 1s 15ms/step - loss: 0.1595 - accuracy: 0.9453 - val_loss: 0.0938 - val_accuracy: 0.9649
Epoch 6/10
64/64 [=====] - 1s 15ms/step - loss: 0.1598 - accuracy: 0.9531 - val_loss: 0.0954 - val_accuracy: 0.9825
Epoch 7/10
64/64 [=====] - 1s 15ms/step - loss: 0.1864 - accuracy: 0.9453 - val_loss: 0.0929 - val_accuracy: 0.9649
Epoch 8/10
64/64 [=====] - 1s 15ms/step - loss: 0.1632 - accuracy: 0.9551 - val_loss: 0.1017 - val_accuracy: 0.9649
Epoch 9/10
64/64 [=====] - 1s 15ms/step - loss: 0.1546 - accuracy: 0.9531 - val_loss: 0.0889 - val_accuracy: 0.9649
Epoch 10/10
64/64 [=====] - 1s 15ms/step - loss: 0.1475 - accuracy: 0.9551 - val_loss: 0.0781 - val_accuracy: 0.9649
*****
*****
Epoch 1/10
2021-12-04 16:58:11.170527: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
2021-12-04 16:58:11.313024: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
7/65 [==>.....] - ETA: 1s - loss: 0.0560 - accuracy: 1.0000
2021-12-04 16:58:11.390411: I tensorflow/core/grappler/optimizers/custom_graph

```

```

h_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
65/65 [=====] - 2s 23ms/step - loss: 0.1538 - accuracy: 0.9552 - val_loss: 0.1336 - val_accuracy: 0.9286
Epoch 2/10
 1/65 [.....] - ETA: 1s - loss: 0.1620 - accuracy: 1.0000
2021-12-04 16:58:12.789349: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
2021-12-04 16:58:12.837932: I tensorflow/core/grappler/optimizers/custom_graph_optimizer_registry.cc:112] Plugin optimizer for device_type GPU is enabled.
65/65 [=====] - 1s 19ms/step - loss: 0.1701 - accuracy: 0.9513 - val_loss: 0.0907 - val_accuracy: 0.9821
Epoch 3/10
65/65 [=====] - 1s 19ms/step - loss: 0.1622 - accuracy: 0.9513 - val_loss: 0.0672 - val_accuracy: 1.0000
Epoch 4/10
65/65 [=====] - 1s 18ms/step - loss: 0.1448 - accuracy: 0.9552 - val_loss: 0.0881 - val_accuracy: 0.9821
Epoch 5/10
65/65 [=====] - 1s 18ms/step - loss: 0.1416 - accuracy: 0.9649 - val_loss: 0.0796 - val_accuracy: 0.9821
Epoch 6/10
65/65 [=====] - 1s 18ms/step - loss: 0.1351 - accuracy: 0.9552 - val_loss: 0.0843 - val_accuracy: 0.9643
Epoch 7/10
65/65 [=====] - 1s 18ms/step - loss: 0.1481 - accuracy: 0.9571 - val_loss: 0.1075 - val_accuracy: 0.9643
Epoch 8/10
65/65 [=====] - 1s 18ms/step - loss: 0.1487 - accuracy: 0.9532 - val_loss: 0.1130 - val_accuracy: 0.9643
Epoch 9/10
65/65 [=====] - 1s 18ms/step - loss: 0.1562 - accuracy: 0.9532 - val_loss: 0.0835 - val_accuracy: 0.9821
Epoch 10/10
65/65 [=====] - 1s 19ms/step - loss: 0.1422 - accuracy: 0.9552 - val_loss: 0.0801 - val_accuracy: 0.9821

```

In [17]:

lstm_metrics_df

Out[17]:

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy |
|---|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|
| 1 | 17.0 | 5.0 | 0.0 | 35.0 | 1.000000 | 0.772727 | 0.875000 | 0.912281 |
| 2 | 23.0 | 2.0 | 6.0 | 26.0 | 0.812500 | 0.920000 | 0.928571 | 0.859649 |
| 3 | 15.0 | 1.0 | 7.0 | 34.0 | 0.829268 | 0.937500 | 0.971429 | 0.859649 |
| 4 | 18.0 | 2.0 | 1.0 | 36.0 | 0.972973 | 0.900000 | 0.947368 | 0.947368 |
| 5 | 17.0 | 1.0 | 6.0 | 33.0 | 0.846154 | 0.944444 | 0.970588 | 0.877193 |
| 6 | 21.0 | 1.0 | 11.0 | 24.0 | 0.685714 | 0.954545 | 0.960000 | 0.789474 |
| 7 | 21.0 | 2.0 | 3.0 | 31.0 | 0.911765 | 0.913043 | 0.939394 | 0.912281 |
| 8 | 22.0 | 1.0 | 5.0 | 29.0 | 0.852941 | 0.956522 | 0.966667 | 0.894737 |
| 9 | 17.0 | 1.0 | 3.0 | 36.0 | 0.923077 | 0.944444 | 0.972973 | 0.929825 |

True Negative False Positive False Negative True Positivity Sensitivity Specificity Precision Accuracy

Cumulative metrics

```
In [29]: all_dfs = [svm_metrics_df, kn_metrics_df, lstm_metrics_df]
all_names = ['SVM', 'KNN', 'LSTM']
```

1st Fold

```
In [34]: # lstm_metrics_df.loc[1,:]
df = pd.DataFrame(columns=performance_metrics)
fold_count = 1
for i, each_df in enumerate(all_dfs):
    temp_df = each_df.xs(fold_count)
    temp_df.name = all_names[i]
    df = df.append(temp_df)
df
```

```
Out[34]:
```

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy | F1 |
|-------------|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|------|
| SVM | 22.0 | 0.0 | 0.0 | 35.0 | 1.0 | 1.000000 | 1.000000 | 1.000000 | 1.0 |
| KNN | 21.0 | 1.0 | 0.0 | 35.0 | 1.0 | 0.954545 | 0.972222 | 0.982456 | 0.98 |
| LSTM | 17.0 | 5.0 | 0.0 | 35.0 | 1.0 | 0.772727 | 0.875000 | 0.912281 | 0.89 |

2nd Fold

```
In [35]: # lstm_metrics_df.loc[1,:]
df = pd.DataFrame(columns=performance_metrics)
fold_count = 2
for i, each_df in enumerate(all_dfs):
    temp_df = each_df.xs(fold_count)
    temp_df.name = all_names[i]
    df = df.append(temp_df)
df
```

```
Out[35]:
```

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy | F1 |
|-------------|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|------|
| SVM | 23.0 | 2.0 | 1.0 | 31.0 | 0.96875 | 0.92 | 0.939394 | 0.947368 | 0.94 |
| KNN | 21.0 | 4.0 | 0.0 | 32.0 | 1.00000 | 0.84 | 0.888889 | 0.929825 | 0.91 |
| LSTM | 23.0 | 2.0 | 6.0 | 26.0 | 0.81250 | 0.92 | 0.928571 | 0.859649 | 0.88 |

3rd Fold

In [36]:

```
# lstm_metrics_df.loc[1,:]
df = pd.DataFrame(columns=performance_metrics)
fold_count = 3
for i, each_df in enumerate(all_dfs):
    temp_df = each_df.xs(fold_count)
    temp_df.name = all_names[i]
    df = df.append(temp_df)
df
```

Out[36]:

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy | F1 |
|-------------|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|------|
| SVM | 15.0 | 1.0 | 1.0 | 40.0 | 0.975610 | 0.9375 | 0.975610 | 0.964912 | 0.95 |
| KNN | 14.0 | 2.0 | 1.0 | 40.0 | 0.975610 | 0.8750 | 0.952381 | 0.947368 | 0.93 |
| LSTM | 15.0 | 1.0 | 7.0 | 34.0 | 0.829268 | 0.9375 | 0.971429 | 0.859649 | 0.88 |

4th Fold

In [37]:

```
# lstm_metrics_df.loc[1,:]
df = pd.DataFrame(columns=performance_metrics)
fold_count = 4
for i, each_df in enumerate(all_dfs):
    temp_df = each_df.xs(fold_count)
    temp_df.name = all_names[i]
    df = df.append(temp_df)
df
```

Out[37]:

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy | F1 |
|-------------|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|------|
| SVM | 20.0 | 0.0 | 0.0 | 37.0 | 1.000000 | 1.00 | 1.000000 | 1.000000 | 1.00 |
| KNN | 19.0 | 1.0 | 0.0 | 37.0 | 1.000000 | 0.95 | 0.973684 | 0.982456 | 0.98 |
| LSTM | 18.0 | 2.0 | 1.0 | 36.0 | 0.972973 | 0.90 | 0.947368 | 0.947368 | 0.94 |

5th Fold

In [38]:

```
# lstm_metrics_df.loc[1,:]
df = pd.DataFrame(columns=performance_metrics)
fold_count = 5
for i, each_df in enumerate(all_dfs):
    temp_df = each_df.xs(fold_count)
    temp_df.name = all_names[i]
    df = df.append(temp_df)
df
```

Out[38]:

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy | F1 |
|------------|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|------|
| SVM | 17.0 | 1.0 | 0.0 | 39.0 | 1.000000 | 0.944444 | 0.975000 | 0.982456 | 0.98 |

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy | F1 |
|--|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|----|
|--|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|----|

6th Fold

```
In [39]: # lstm_metrics_df.loc[1,:]
df = pd.DataFrame(columns=performance_metrics)
fold_count = 6
for i, each_df in enumerate(all_dfs):
    temp_df = each_df.xs(fold_count)
    temp_df.name = all_names[i]
    df = df.append(temp_df)
df
```

```
Out[39]:
```

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy | F1 |
|-------------|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|-----|
| SVM | 19.0 | 3.0 | 0.0 | 35.0 | 1.000000 | 0.863636 | 0.921053 | 0.947368 | 0.9 |
| KNN | 19.0 | 3.0 | 0.0 | 35.0 | 1.000000 | 0.863636 | 0.921053 | 0.947368 | 0.9 |
| LSTM | 21.0 | 1.0 | 11.0 | 24.0 | 0.685714 | 0.954545 | 0.960000 | 0.789474 | 0.8 |

7th Fold

```
In [40]: # lstm_metrics_df.loc[1,:]
df = pd.DataFrame(columns=performance_metrics)
fold_count = 7
for i, each_df in enumerate(all_dfs):
    temp_df = each_df.xs(fold_count)
    temp_df.name = all_names[i]
    df = df.append(temp_df)
df
```

```
Out[40]:
```

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy | F1 |
|-------------|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|-----|
| SVM | 22.0 | 1.0 | 1.0 | 33.0 | 0.970588 | 0.956522 | 0.970588 | 0.964912 | 0.9 |
| KNN | 22.0 | 1.0 | 0.0 | 34.0 | 1.000000 | 0.956522 | 0.971429 | 0.982456 | 0.9 |
| LSTM | 21.0 | 2.0 | 3.0 | 31.0 | 0.911765 | 0.913043 | 0.939394 | 0.912281 | 0.9 |

8th Fold

In [42]:

```
# lstm_metrics_df.loc[1,:]
df = pd.DataFrame(columns=performance_metrics)
fold_count = 8
for i, each_df in enumerate(all_dfs):
    temp_df = each_df.xs(fold_count)
    temp_df.name = all_names[i]
    df = df.append(temp_df)
df
```

Out[42]:

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy | F1 |
|-------------|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|----------|
| SVM | 23.0 | 0.0 | 2.0 | 32.0 | 0.941176 | 1.000000 | 1.000000 | 0.964912 | 0.972603 |
| KNN | 23.0 | 0.0 | 0.0 | 34.0 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |
| LSTM | 22.0 | 1.0 | 5.0 | 29.0 | 0.852941 | 0.956522 | 0.966667 | 0.894737 | 0.905660 |

9th Fold

In [43]:

```
# lstm_metrics_df.loc[1,:]
df = pd.DataFrame(columns=performance_metrics)
fold_count = 9
for i, each_df in enumerate(all_dfs):
    temp_df = each_df.xs(fold_count)
    temp_df.name = all_names[i]
    df = df.append(temp_df)
df
```

Out[43]:

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy | F1 |
|-------------|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|----------|
| SVM | 18.0 | 0.0 | 0.0 | 39.0 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |
| KNN | 18.0 | 0.0 | 0.0 | 39.0 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |
| LSTM | 17.0 | 1.0 | 3.0 | 36.0 | 0.923077 | 0.944444 | 0.972973 | 0.929825 | 0.934274 |

10th Fold

In [44]:

```
# lstm_metrics_df.loc[1,:]
df = pd.DataFrame(columns=performance_metrics)
fold_count = 10
for i, each_df in enumerate(all_dfs):
    temp_df = each_df.xs(fold_count)
    temp_df.name = all_names[i]
    df = df.append(temp_df)
df
```

Out[44]:

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy | F1 |
|------------|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|----------|
| SVM | 25.0 | 0.0 | 0.0 | 31.0 | 1.00000 | 1.00 | 1.000000 | 1.000000 | 1.000000 |

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy | F1 |
|--|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|----|
|--|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|----|

Average of all

```
In [45]: # lstm_metrics_df.loc[1,:]
df = pd.DataFrame(columns=performance_metrics)
fold_count = 'Average'
for i, each_df in enumerate(all_dfs):
    temp_df = each_df.xs(fold_count)
    temp_df.name = all_names[i]
    df = df.append(temp_df)
df
```

```
Out[45]:
```

| | True Negative | False Positive | False Negative | True Positivity | Sensitivity | Specificity | Precision | Accuracy | F1 |
|-------------|------------------|-------------------|-------------------|--------------------|-------------|-------------|-----------|----------|-----|
| SVM | 20.4 | 0.8 | 0.5 | 35.2 | 0.985612 | 0.962210 | 0.978164 | 0.977193 | 0.9 |
| KNN | 19.7 | 1.5 | 0.2 | 35.5 | 0.994997 | 0.930415 | 0.959341 | 0.970113 | 0.9 |
| LSTM | 19.6 | 1.6 | 4.7 | 31.0 | 0.867310 | 0.924323 | 0.953199 | 0.889317 | 0.9 |

Observation

- I consider balanced accuracy to be the optimal metric to find the best model.
- The case being, SVM is the model which is giving the highest balanced accuracy.

Why is SVM performing better?

- SVM doesn't get affected by outliers
- It does not suffer from overfitting
- It is more efficient than other ML algorithms listed here